

Machine Learning Clusterização

Agrupamento de Clientes Por Consumo de Energia

A partir de dados de consumo de energia de clientes, vamos agrupar os consumidores por similaridade a fim de compreender o comportamento dos clientes e sua relação com o consumo de energia.

```
In [83]: # Imports
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from matplotlib import pylab
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split
from scipy.spatial.distance import cdist, pdist
from sklearn.metrics import silhouette_score
import warnings
warnings.filterwarnings("ignore")
%matplotlib inline
```

```
In [84]: # Carregando os dados
dataset = pd.read_csv('consumo_energia.txt', delimiter = ';')
```

```
In [85]: # Visualizando informações sobre o dataset
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2075259 entries, 0 to 2075258
Data columns (total 9 columns):
 #   Column                Dtype
---  -
 0   Date                  object
 1   Time                  object
 2   Global_active_power   object
 3   Global_reactive_power object
 4   Voltage               object
 5   Global_intensity      object
 6   Sub_metering_1        object
 7   Sub_metering_2        object
 8   Sub_metering_3        float64
dtypes: float64(1), object(8)
memory usage: 142.5+ MB
```

```
In [86]: dataset.head()
```

Out[86]:

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub
0	16/12/2006	17:24:00	4.216	0.418	234.840	18.400	
1	16/12/2006	17:25:00	5.360	0.436	233.630	23.000	
2	16/12/2006	17:26:00	5.374	0.498	233.290	23.000	
3	16/12/2006	17:27:00	5.388	0.502	233.740	23.000	
4	16/12/2006	17:28:00	3.666	0.528	235.680	15.800	

In [87]: `dataset.shape`

Out[87]: (2075259, 9)

In [88]: `dataset.dtypes`

Out[88]:

Date	object
Time	object
Global_active_power	object
Global_reactive_power	object
Voltage	object
Global_intensity	object
Sub_metering_1	object
Sub_metering_2	object
Sub_metering_3	float64
dtype:	object

In [89]: `# Checando se há valores missing`
`dataset.isnull().values.any()`

Out[89]: True

In [90]: `# Remove os registros com valores NA e remove as duas primeiras colunas (não são ne`
`dataset = dataset.iloc[0:, 2:9].dropna()`

In [91]: `dataset['Voltage'] = dataset['Voltage'].astype(dtype = 'float64')`
`dataset['Global_active_power'] = dataset['Global_active_power'].astype(dtype = 'flo`
`dataset['Global_reactive_power'] = dataset['Global_reactive_power'].astype(dtype =`
`dataset['Global_intensity'] = dataset['Global_intensity'].astype(dtype = 'float64')`
`dataset['Sub_metering_1'] = dataset['Sub_metering_1'].astype(dtype = 'float64')`
`dataset['Sub_metering_2'] = dataset['Sub_metering_2'].astype(dtype = 'float64')`
`dataset['Sub_metering_3'] = dataset['Sub_metering_3'].astype(dtype = 'float64')`

In [92]: `dataset.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2049280 entries, 0 to 2075258
Data columns (total 7 columns):
 #   Column                Dtype
---  -
 0   Global_active_power    float64
 1   Global_reactive_power  float64
 2   Voltage                float64
 3   Global_intensity       float64
 4   Sub_metering_1         float64
 5   Sub_metering_2         float64
 6   Sub_metering_3         float64
dtypes: float64(7)
memory usage: 125.1 MB
```

In [93]: `dataset.head()`

```
Out[93]:
```

	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_met
0	4.216	0.418	234.84	18.4	0.0	
1	5.360	0.436	233.63	23.0	0.0	
2	5.374	0.498	233.29	23.0	0.0	
3	5.388	0.502	233.74	23.0	0.0	
4	3.666	0.528	235.68	15.8	0.0	

In [94]: `# Checando se há valores missing`
`dataset.isnull().values.any()`

Out[94]: False

In [95]: `# Obtém os valores dos atributos. Neste caso as variaveis foram carregadas como cat`
`dataset_atrib = dataset.values`

In [96]: `dataset_atrib`

```
Out[96]: array([[ 4.216,  0.418, 234.84, ...,  0.    ,  1.    , 17.    ],
        [ 5.36 ,  0.436, 233.63, ...,  0.    ,  1.    , 16.    ],
        [ 5.374,  0.498, 233.29, ...,  0.    ,  2.    , 17.    ],
        ...,
        [ 0.938,  0.    , 239.82, ...,  0.    ,  0.    ,  0.    ],
        [ 0.934,  0.    , 239.7 , ...,  0.    ,  0.    ,  0.    ],
        [ 0.932,  0.    , 239.55, ...,  0.    ,  0.    ,  0.    ]])
```

In [97]: `# Coleta uma amostra de 1% dos dados para não comprometer a memória do computador`
`amostra1, amostra2 = train_test_split(dataset_atrib, train_size = .01)`

In [98]: `amostra1.shape`

Out[98]: (20492, 7)

In [99]: `# Aplica redução de dimensionalidade`
`# Transforma as 7 variáveis em 2 variaveis principais. Esse método utiliza Algebra`
`# entre os dados e assim "juntar" as variaveis, medindo a semelhança pela variância`
`pca = PCA(n_components = 2).fit_transform(amostra1)`

In [100... `# Determinando um range de K`
`k_range = range(1,12)`

In [103... `# Aplicando o modelo K-Means para cada valor de K (esta célula pode levar bastante`
`k_means_var = [KMeans(n_clusters = k).fit(pca) for k in k_range]`

In [104... `# Ajustando o centróide do cluster para cada modelo`
`centroids = [X.cluster_centers_ for X in k_means_var]`

In [105... `# Calculando a distância euclidiana de cada ponto de dado para o centróide`
`k_euclid = [cdist(pca, cent, 'euclidean') for cent in centroids]`
`dist = [np.min(ke, axis = 1) for ke in k_euclid]`

In [106... `# Soma dos quadrados das distâncias dentro do cluster`
`soma_quadrados_intra_cluster = [sum(d**2) for d in dist]`

```


In [107... # Soma total dos quadrados
soma_total = sum(pdist(pca)**2)/pca.shape[0]

In [108... # Soma dos quadrados entre clusters
soma_quadrados_inter_cluster = soma_total - soma_quadrados_intra_cluster

In [109... # Curva de Elbow
fig = plt.figure()
ax = fig.add_subplot(111)
ax.plot(k_range, soma_quadrados_inter_cluster/soma_total * 100, 'b*-')
ax.set_ylim((0,100))
plt.grid(True)
plt.xlabel('Número de Clusters')
plt.ylabel('Percentual de Variância Explicada')
plt.title('Variância Explicada x Valor de K')

```

Out[109]: Text(0.5, 1.0, 'Variância Explicada x Valor de K')



```

In [112... # Criando um modelo com K = 8
modelo_v1 = KMeans(n_clusters = 8)
modelo_v1.fit(pca)

```

Out[112]: KMeans()

In []:

In []:

```

In [113... # Obtém os valores mínimos e máximos e organiza o shape
x_min, x_max = pca[:, 0].min() - 5, pca[:, 0].max() + 1
y_min, y_max = pca[:, 1].min() + 1, pca[:, 1].max() + 5
xx, yy = np.meshgrid(np.arange(x_min, x_max, .02), np.arange(y_min, y_max, .02))

```


```
Z = modelo_v1.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)
```

In [114...]

```
# Plot das reas dos clusters
plt.figure(1)
plt.clf()
plt.imshow(Z,
            interpolation = 'nearest',
            extent = (xx.min(), xx.max(), yy.min(), yy.max()),
            cmap = plt.cm.Paired,
            aspect = 'auto',
            origin = 'lower')
```


Out[114]:

```
<matplotlib.image.AxesImage at 0x23c8588cd30>
```



In [115...]

```
# Plot dos centrides
plt.plot(pca[:, 0], pca[:, 1], 'k.', markersize = 4)
centroids = modelo_v1.cluster_centers_
inert = modelo_v1.inertia_
plt.scatter(centroids[:, 0], centroids[:, 1], marker = 'x', s = 169, linewidths = 3)
plt.xlim(x_min, x_max)
plt.ylim(y_min, y_max)
plt.xticks(())
plt.yticks(())
plt.show()
```



In [117... `#?silhouette_score`

In [118... `# Silhouette Score`
`labels = modelo_v1.labels_`
`silhouette_score(pca, labels, metric = 'euclidean')`

Out[118]: 0.8088812305122378


In [119... `# Criando um modelo com K = 10`
`modelo_v2 = KMeans(n_clusters = 10)`
`modelo_v2.fit(pca)`

Out[119]: KMeans(n_clusters=10)


In [120... `# Obtém os valores mínimos e máximos e organiza o shape`
`x_min, x_max = pca[:, 0].min() - 5, pca[:, 0].max() + 1`
`y_min, y_max = pca[:, 1].min() + 1, pca[:, 1].max() + 5`
`xx, yy = np.meshgrid(np.arange(x_min, x_max, .02), np.arange(y_min, y_max, .02))`
`Z = modelo_v2.predict(np.c_[xx.ravel(), yy.ravel()])`
`Z = Z.reshape(xx.shape)`

In [121... `# Plot das áreas dos clusters`
`plt.figure(1)`
`plt.clf()`
`plt.imshow(Z,`
`interpolation = 'nearest',`
`extent = (xx.min(), xx.max(), yy.min(), yy.max()),`
`cmap = plt.cm.Paired,`
`aspect = 'auto',`
`origin = 'lower')`

Out[121]: <matplotlib.image.AxesImage at 0x23c998d0040>



```
In [122... # Plot dos centróides
plt.plot(pca[:, 0], pca[:, 1], 'k.', markersize = 4)
centroids = modelo_v2.cluster_centers_
inert = modelo_v2.inertia_
plt.scatter(centroids[:, 0], centroids[:, 1], marker = 'x', s = 169, linewidths = 3)
plt.xlim(x_min, x_max)
plt.ylim(y_min, y_max)
plt.xticks(())
plt.yticks(())
plt.show()
```



```
In [123... # Silhouette Score -- utilizamos para avaliar se o numero de cluster é o ideal... n
labels = modelo_v2.labels_
silhouette_score(pca, labels, metric = 'euclidean')
```

```
Out[123]: 0.682286610628249
```

Criando o Cluster Map com os clusters do Modelo V1 que apresentou melhor Silhouette Score.

```
In [124... # Lista com nomes das colunas
names = ['Global_active_power', 'Global_reactive_power', 'Voltage', 'Global_intensi
```

```
In [125... # Cria o cluster map
cluster_map = pd.DataFrame(amostra1, columns = names)
cluster_map['Global_active_power'] = pd.to_numeric(cluster_map['Global_active_power']
cluster_map['cluster'] = modelo_v1.labels_
```

```
In [128... cluster_map
```

```
Out[128]:
```

	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub
0	0.318	0.114	242.34	1.4	0.0	
1	2.184	0.134	242.47	9.0	0.0	
2	0.392	0.366	243.91	2.2	0.0	
3	3.426	0.000	232.05	14.6	35.0	
4	1.514	0.194	239.18	6.4	0.0	
...
20487	0.394	0.000	241.34	1.6	2.0	
20488	1.450	0.108	240.64	6.0	0.0	
20489	0.320	0.084	242.66	1.4	0.0	
20490	0.376	0.000	243.32	1.6	0.0	
20491	0.202	0.000	239.46	0.8	0.0	

20492 rows × 8 columns

```
In [129... # Calcula a média de consumo de energia por cluster
cluster_map.groupby('cluster')['Global_active_power'].mean()
```

```
Out[129]:
```

cluster	Global_active_power
0	0.515455
1	1.810525
2	4.635189
3	3.511884
4	3.776012
5	1.106406
6	2.352873
7	2.556937

Name: Global_active_power, dtype: float64

```
In [130... # Calcula a quantidade de observacoes por cluster
cluster_map.groupby('cluster')['Global_active_power'].count()
```



```
Out[130]: cluster
0      12853
1       5894
2        397
3        207
4        324
5        399
6        181
7        237
Name: Global_active_power, dtype: int64
```

```
In [ ]:
```

Gerando Cluster com dados normalizados

```
In [131... # Obtém os valores dos atributos. Neste caso as variaveis foram carregadas como cat
dataset_atrib = dataset.values

# Importa biblioteca para fazer a normalizacao
from sklearn.preprocessing import MinMaxScaler

# Cria o objeto para normalizar e faz a normalizacao dos dados
Padronizador = MinMaxScaler()
dataset_atrib = Padronizador.fit_transform(dataset_atrib)

amostra1, amostra2 = train_test_split(dataset_atrib, train_size = .01)
pca = PCA(n_components = 2).fit_transform(amostra1)
k_range = range(1,12)
k_means_var = [KMeans(n_clusters = k).fit(pca) for k in k_range]
centroids = [X.cluster_centers_ for X in k_means_var]
k_euclid = [cdist(pca, cent, 'euclidean') for cent in centroids]
dist = [np.min(ke, axis = 1) for ke in k_euclid]
soma_quadrados_intra_cluster = [sum(d**2) for d in dist]
soma_total = sum(pdist(pca)**2)/pca.shape[0]
# Soma dos quadrados entre clusters
soma_quadrados_inter_cluster = soma_total - soma_quadrados_intra_cluster
# Criando um modelo com K = 8
modelo_v1 = KMeans(n_clusters = 8)
modelo_v1.fit(pca)
# Obtém os valores mínimos e máximos e organiza o shape
x_min, x_max = pca[:, 0].min() - 5, pca[:, 0].max() - 1
y_min, y_max = pca[:, 1].min() + 1, pca[:, 1].max() + 5
xx, yy = np.meshgrid(np.arange(x_min, x_max, .02), np.arange(y_min, y_max, .02))
Z = modelo_v1.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)
# Silhouette Score
labels = modelo_v1.labels_
silhouette_score(pca, labels, metric = 'euclidean')
# Lista com nomes das colunas
names = ['Global_active_power', 'Global_reactive_power', 'Voltage', 'Global_intensi

# Cria o cluster map
cluster_map = pd.DataFrame(amostra1, columns = names)
cluster_map['Global_active_power'] = pd.to_numeric(cluster_map['Global_active_power
cluster_map['cluster'] = modelo_v1.labels_
cluster_map
```

Out[131]:

	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Su
0	0.113706	0.000000	0.619063	0.107884	0.011364	
1	0.019736	0.000000	0.491438	0.029046	0.000000	
2	0.016658	0.000000	0.745719	0.016598	0.000000	
3	0.015571	0.113669	0.665267	0.020747	0.000000	
4	0.167663	0.034532	0.470436	0.161826	0.000000	
...
20487	0.017563	0.054676	0.589338	0.020747	0.000000	
20488	0.055948	0.166906	0.680775	0.058091	0.000000	
20489	0.081115	0.195683	0.602585	0.082988	0.000000	
20490	0.033134	0.143885	0.586430	0.037344	0.000000	
20491	0.138874	0.000000	0.487561	0.132780	0.011364	

20492 rows × 8 columns



In [132...]

```
# Calcula da quantidade de observações por cluster  
cluster_map.groupby('cluster')['Global_active_power'].count()
```

Out[132]:

```
cluster  
0      4276  
1      3639  
2      9073  
3       627  
4      1750  
5       536  
6       417  
7       174  
Name: Global_active_power, dtype: int64
```

In []:

In []: